

1. We adopt the squared error loss:  $L = \frac{1}{2} (h(x_1, x_2) - y)^2$

$$\text{Let } z = b + w_1 x_1 + w_2 x_2 \Rightarrow h(x_1, x_2) = \sigma(z)$$

Since  $\sigma'(z) = \sigma(z)(1 - \sigma(z))$ , the gradients with respect to the parameters are:

$$\text{For } b: \frac{\partial L}{\partial b} = (h - y) \sigma(z) (1 - \sigma(z))$$

$$\text{For } w_1: \frac{\partial L}{\partial w_1} = (h - y) \sigma(z) (1 - \sigma(z)) \cdot x_1$$

$$\text{For } w_2: \frac{\partial L}{\partial w_2} = (h - y) \sigma(z) (1 - \sigma(z)) \cdot x_2$$

Thus, the gradient vector is  $\nabla_{\theta} L = (h - y) \sigma(z) (1 - \sigma(z)) \cdot (1, x_1, x_2)$

The standard SGD update is  $\theta' = \theta - \eta \nabla_{\theta} L$ , where  $\eta$  is the learning rate

For the given data point and parameter initialization, we compute

$$z = 4 + 5(1) + 6(2) = 21, \quad h = \sigma(21)$$

Hence, the update becomes  $\theta' = (4, 5, 6) - \eta \cdot ((\sigma(21) - 3) \sigma(21) (1 - \sigma(21))) \cdot (1, 1, 2)$

Explicitly, this gives:  $b' = 4 - \eta \cdot ((\sigma(21) - 3) \sigma(21) (1 - \sigma(21)))$

$$w_1' = 5 - \eta \cdot ((\sigma(21) - 3) \sigma(21) (1 - \sigma(21))) \cdot 1$$

$$w_2' = 6 - \eta \cdot ((\sigma(21) - 3) \sigma(21) (1 - \sigma(21))) \cdot 2$$

2. (a) The sigmoid function is  $\sigma(x) = \frac{1}{1+e^{-x}}$

For  $k=1$ ,  $\sigma'(x) = \sigma(x)(1-\sigma(x))$

For  $k=2$ ,  $\sigma''(x) = \sigma'(x)(1-2\sigma(x)) = \sigma(x)(1-\sigma(x))(1-2\sigma(x))$

For  $k=3$ ,  $\sigma'''(x) = \sigma(x)(1-\sigma(x))(1-6\sigma(x)+6\sigma(x)^2)$

(b) The sigmoid function is closely related to the hyperbolic tangent function. In fact, we have:  $\sigma(x) = \frac{1}{1+e^{-x}} = \frac{1}{2} (1 + \tanh(\frac{x}{2}))$ .

This identity is useful because it connects the logistic function to hyperbolic functions, which often arise in the analysis of differential equations and neural network activations.

3. Alternative activation functions: Why do modern neural networks often prefer the Relu or other nonlinearities instead of the sigmoid function?

Gradient vanishing problem: Since the sigmoid function saturates at very large positive or negative inputs, how does this affect gradient propagation in deep networks? What are common solutions?