

Approximating the Runge Function with a Small Neural Network

➤ Target

$$f(x) = \frac{1}{1 + 25x^2}, \quad x \in [-1, 1]$$

➤ Method

Model One-hidden-layer MLP (width H=64) with a leaky-ReLU hidden activation and linear output.

Training objective Mean-squared error (MSE) between prediction $\hat{f}(x)$ and $f(x)$.

Data 256 training points on a uniform grid in $[-1, 1]$; validation uses 256 **midpoints** to avoid overlap.

Optimization Full-batch gradient descent, learning rate 0.010.010.01 with step decays; 3000 epochs.

Implementation Pure-Python (no external libraries), forward & back-prop coded by hand.

➤ Loss definitions

Function loss (MSE)

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (\hat{f}(x_i) - f(x_i))^2$$

Max error (on a dense grid)

$$\max_i |\hat{f}(x_i) - f(x_i)|$$

➤ Results

Learning curves (numbers from your logs)

Epoch	Train MSE	Val MSE
200	0.017319	0.017174
1000	0.010072	0.010022
2000	0.003475	0.003463
3000	0.001768	0.001763

Both curves decrease steadily and remain almost identical → no overfitting.

Final evaluation (dense grid of 1000 points)

MSE: 0.00176431

Max error: 0.07533909 at $x \approx 0.0771$

These errors are small relative to the function scale $[0.0385, 1]$, showing a solid fit.

Qualitative behavior

The network captures the central peak near $x=0$ and the flat tails toward ± 1 .

The largest deviation occurs around a moderate x (here $x \approx 0.077$), which matches where the curvature changes fastest.

➤ Discussion

What helped

Using midpoints for validation avoids train/val leakage on a grid.

A simple learning-rate decay after plateaus.

Even a single hidden layer is sufficient for this 1D target.

Limitations / possible improvements

Increase width (e.g., $H=128$) or train longer to push MSE lower.

Use Chebyshev-biased sampling (denser near ± 1) to reduce edge errors for Runge-type shapes.

A smooth activation (tanh-like) can further stabilize training if you later add a derivative term to the loss.